

---

# Optical Flow Regularization of Implicit Neural Representations for Video Frame Interpolation

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Recent works have shown the ability of neural implicit representations (NIR) to carry meaningful representations of signal derivatives. In this work, we leverage this property to perform video frame interpolation by explicitly constraining the derivatives of the NIR to satisfy the optical flow constraint equation. We achieve state of the art video frame interpolation on limited motion ranges using only a target video and its optical flow, without learning the interpolation operator from additional training data. We further show that constraining the NIR derivatives not only allows to interpolate intermediate frames but also improves the ability of narrow networks to fit observed frames, which suggests potential applications to NIR optimization and video compression.

## 1 Introduction

Many core concepts across the fields of signal processing are defined in terms of continuous functions and their derivatives: surfaces are continuous manifolds in space, motion is a rate of change in space through time, etc. In contrast, the modern digital infrastructure is inherently discrete: digital sensors capture discrete observations of the world sampled in time and space; digital computers store and process discrete representations of signals. In order to model continuous notions on discrete signal representations, classical signal processing approaches have resorted to a variety of heuristics and assumptions, often taking the form of constant first or second derivatives of the signal between consecutive observations. The lack of generality of any such handcrafted heuristics, combined with the ever improving quantitative results of Machine Learning (ML) approaches, have led to the near ubiquitous use of ML approaches in recent signal processing research. These approaches leverage large collections of data to infer statistical properties of signals instead of hand-crafted heuristics.

In computer vision, Video Frame Interpolation (VFI) is one task representative of such development. VFI models aim to infer intermediate frames between consecutive frames of a video. To do so, most successful approaches rely on the optical flow as an approximation of the motion field to guide the interpolation of pixel intensities from the grid of two consecutive frames onto the pixel grid of intermediate frames. Classical approaches formulate assumptions such as constant speed or acceleration of the motion field between consecutive frames [CITE]. The value of each pixel in the inferred intermediate frame is computed by first shifting the pixel intensities of the observed frames following the optical flow directions, and then interpolating the shifted pixel intensities onto the intermediate frame's pixel grid. These approaches suffer from the following two limitations:

- Optical flow constraint used to infer the optical flow holds for limited situations.
- Linear interpolation of pixel intensities along the optical flow directions does not hold in practice.

These limitations share a common root cause: discretization. Indeed, both the optical flow constraint and the constant motion field assumption only truly hold at the infinitesimal scale, for much smaller time deltas than typical FPS used in practice.

ML approaches [CITE] have instead proposed to learn the frame interpolation operator from large video collections, without explicitly formulating any assumption on the optical flow. While these approaches have achieved great success in terms of benchmark performance, they are prone to generalization errors when applied to unseen videos. Indeed differences between the training set distribution (i.e. VFI benchmark videos) and the target video distribution hinders the performance of ML approaches: differences in the range of motion, exposure time and frame-per-second have been shown to limit the generalization of state-of-the-art models to video frame interpolation in the wild [CITE].

In the mean time, research on implicit representations seek better discrete representations of continuous signals. In recent years Neural Implicit Representations (NIR), i.e. representing signals as Neural Networks (NN) have been shown to offer several competitive advantages over explicit representations, with notable early successes for 3D shape representations [CITE]. Of particular interest to us is the work of SIREN [CITE], in which it has been shown that representing signals using Multi Layer Perceptrons (MLP) with sine activation functions carry meaningful representations of the signal derivatives. Inspired by this work, we question whether such approach may be used to guide the interpolation process of VFI by controlling the exact derivatives of the signal rather than finite differences, thus avoiding the discretization pitfalls of traditional approaches. We do so by constraining the derivatives of SIREN representations to satisfy the optical flow constraint, i.e., to be orthogonal to the video’s optical flow (which we compute using existing state-of-the-art OF models). We find that this approach outperforms most existing machine learning-based approaches on small motion range benchmarks, without relying on machine learning for the interpolation operator: we simply regularize the implicit representation to satisfy the definition of the optical flow. In this sense, our approach is most similar to classical VFI approaches, except that instead of wrapping the OF on discrete explicit frame representations, we apply the optical flow constraint on the exact gradient of the the NIR. Our method is thus not subject to any mismatch between training and test data. Furthermore, our approach can sample any number of frame in-between the observed frames due to the continuous nature of the representation. In addition to its application to VFI, we also show that constraining the gradient of the model also improves the ability of narrow MLPs to fit the signal, suggesting potential applications in NIR optimization and video compression.

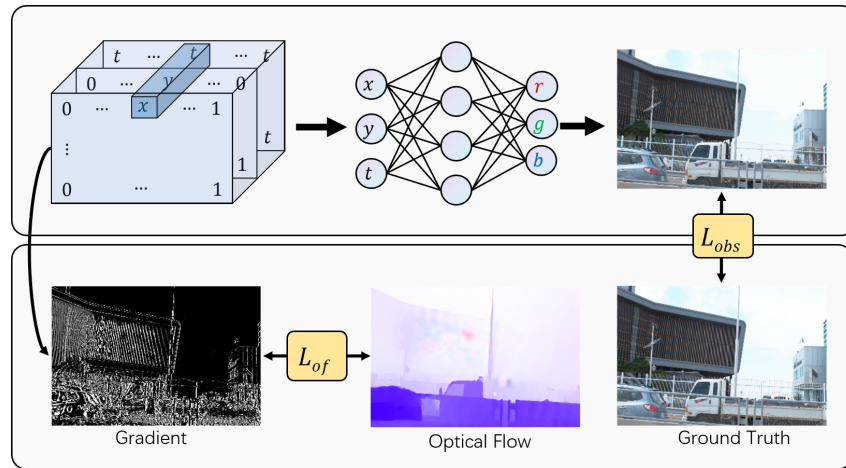


Figure 1: Illustration of our approach

To summarize, the contributions of this work are:

- We propose a regularization method for SIREN which achieve state-of-the-art video frame interpolation on small motion ranges.
- In contrast to other state-of-the-art approaches, our approach does not rely on training on a large external training set. It only relies on the target video and its estimated optical flow.

- We show that our regularization approach not only helps generalizing to intermediate frame generalization but also helps narrow models fit the observed frames.

On the other hand, our approach (in its current form) presents important limitations:

- It relies on an input optical flow, which is computed using existing ML-based model and thus suffers the limitations of ML approaches.
- Optimization of the NIR is very time-consuming, which hinders our ability to work on full resolution videos for time constraints.
- Our method currently only works on limited motion range. It does not match state-of-the-art ML models on large motion ranges.

While we acknowledge the importance of the above limitations, we believe these to not be fundamental limitations of our approach but rather important future NIR research directions. We discuss these limitations at length and present possible axis to tackle them in Section XXX. The remainder of this paper is organized as follows: We briefly present some related work in Section XXX, the detail of our method in Section XXX, and design several experiments to highlight the advantages of our approach in Section XXX.

## 2 Related Work

**Deep learning video interpolation.** A number of deep learning models have been developed for video interpolation tasks. Almost all models can be categorized as: optical flow based, and kernel based.

*Optical Flow-Based.* Optical flow-based approaches are the most popular in video frame interpolation. The standard technique of video frame interpolation aims at explicitly estimating motion in the form of optical flow, warping two input frames to an intermediate frame, and synthesizing the occlusion region. The frames are constrained by the assumption of linear motion and constant luminance between them. However, video interpolation of video frames is heavily dependent on the accuracy of optical flow.

The Super-SloMo ? proposed by Jiang et.al. is a non-negligible work in the task of optical flow-based video frame interpolation. Super-SloMo extends the U-Net architecture proposed by Liu et al ?. The bilateral optical flow is calculated for the input two frames and approximates the key frame with the intermediate optical flow of the two frames. Then the frames of the input are warped according to the obtained intermediate optical flow.

RRIN ? mentioned that the estimation of intermediate frames in Super-SloMo works poorly near the boundaries because the optical flow is not locally smooth in these regions. RRIN proposes to improve the accuracy of optical flow by residual learning. BMBC ? adds two additional approximate vectors to Super-SloMo to make the bilateral motion estimation more accurate.

*Kernel-Based.* To avoid explicit motion estimation and warping stages, the kernel-based approach performs a convolution operation on the input frames and the output of the convolution is used as the result of interpolating the frames. Niklaus et al. ? proposed a fully convolutional deep neural network using a spatially adaptive convolutional kernel to perform the prediction of intermediate frames for two frames with consecutive inputs. Niklaus et al. ? improved their method by using a separable convolution with spatially adaptive one-dimensional convolutional kernel pairs estimated for each pixel, in reducing the parameters of the model. The results of kernel-based methods for frame interpolation can be limited by the size of the kernel.

Lee et al. proposed Adacof ?, which can use any pixel at any position for convolution operation, so that the convolution kernel is no longer limited to the local range. And many methods residing in optical flow are defined as a special case of Adacof. However, most kernel-based methods can only generate one intermediate frame, and if one wants to generate multiple intermediate frames, one needs to do it recursively. EDSC ? is the first kernel-based method proposed to generate multiple intermediate frames, but the results are not as good as the optical flow method.

**Implicit Neural Network Representation. (INR)**

121 INR use a neural network to represent an object approximately, which is essentially a way to  
 122 parameterize the signal. Since ?, ? was developed, INR has performed well in the areas of 3D vision  
 123 tasks, images, and video. The image and video tasks most relevant to this paper are around the  
 124 direction of image/video compression.

125 COIN ? first proposed the use of INR to compress images, mapping pixel coordinates to RGB values.  
 126 COIN++ ? cooperated with the meta-learning approach for image compression work based on COIN.  
 127 In the field of video compression, NeRV ? proposed by Chen et al. successfully encodes the video  
 128 into a neural network, i.e., the content of the video is saved using a neural network. Only the frame  
 129 index of the model needs to be provided, and the corresponding RGB picture is output. In other  
 130 words, this makes it possible to output infinite frames of video using a neural network. Although  
 131 NeRV briefly attempts the task of performing video frame interpolation, this is not NeRV’s main  
 132 work. The NRFF ? proposed by Rho et al., which uses optical flow and residuals information for  
 133 video compression, does not directly fit all frames.

134 Most related to our approach is the concurrent work by XX et al. ?, which also uses INR for video  
 135 interpolation tasks. Their approach, CURE, uses machine learning. It requires visual features of the  
 136 video and does not fully map the pixel coordinates and frame positions of the video to RGB images.

### 137 3 Method

138 We consider a ground-truth video as a continuous signal  $v$  mapping continuous spatial  $(x, y)$  and  
 139 temporal  $(t)$  coordinates to RGB values:

$$\begin{aligned} v : (x, y, t) &\rightarrow (R, G, B) \\ v : \mathbb{R}^3 &\rightarrow \mathbb{R}^3 \end{aligned} \quad (1)$$

140 Our goal is to find a continuous function  $f_\theta$ , parameterized by a finite parameter set  $\theta \in \Theta$ , with  
 141 minimum distance  $d$  to the ground-truth signal:

$$\begin{aligned} f_\theta : (x, y, t) &\rightarrow (R, G, B) \\ s.t. \theta = \min_{\Theta} \iiint d(f_\theta(x, y, t), v(x, y, t)) dx dy dt \end{aligned} \quad (2)$$

142 where the distance function  $d$  may either be the Peak Signal to Noise Ratio (PSNR) or the Structural  
 143 Similarity Index Measure (SSIM). To do so, we only have access to regularly sampled observation of  
 144 the signal  $v$  (i.e. the explicit representation of the video), which we denote as:

$$\begin{aligned} \mathcal{V} &\in \mathbb{R}^{T \times H \times W \times 3} \\ s.t. \mathcal{V}_{xyt} &= v(x, y, t) \end{aligned} \quad (3)$$

145 where  $T$  represents the number of frames in the video, and  $H \times W$  the spatial resolution. We use  
 146 SIREN as parameterized function class  $f_\theta$ . The most straightforward way to approximate Equation 2  
 147 is to optimize the model parameters so as to fit the video frames, using the following loss function we  
 148 refer to as the observation loss:

$$\mathcal{L}_{obs} = \frac{1}{HWT} \sum_{x=1}^W \sum_{y=1}^H \sum_{t=1}^T \|f_\theta(x, y, t) - \mathcal{V}_{xyt}\|^2 \quad (4)$$

149 However, we found that optimizing the NIR to only minimize this observation loss leads to overfitting  
 150 the observation with high temporal frequencies: the intra-frame signal, which we aim to correctly  
 151 recover, shows important deviations from the observed frames, as illustrated in Figure XXX. This  
 152 observation has lead us to consider fitting not only the signal itself, but to also constrain its derivatives.  
 153 In particular, we regularize the model so as to respect the optical flow constraint.

154 The optical flow constraint equation states that for an infinitesimal lapse of time  $\delta t$ , the brightness of  
 155 physical points perceived by a camera at arbitrary coordinates  $(x, y, t)$  should remain constant. In




Figure 2: Illustration of NIR frame interpolation with and without optical flow regularization. Without regularization (middle top), intermediate frames show unnatural high-frequency variations. Regularizing the NIR to satisfy the optical flow constraint equation result in nicely interpolated frames (middle bottom).

other words, given the displacement  $(\delta x, \delta y)$  of a physical point in the image coordinate system, the image brightness  $v$  should remain constant:

$$v(x, y, t) = v(x + \delta x, y + \delta y, t + \delta t) \quad (5)$$

Expressing movement as a ratio of displacement in time and abbreviating coordinates as  $x = (x, y, t)$ , we can write the optical flow  $F$  and the above constraint as:

$$F(x) = \left( \frac{\delta x}{\delta t}, \frac{\delta y}{\delta t}, 1 \right) \quad (6)$$

$$v(x) = v(x + F(x))$$

We leverage this optical flow constraint equation to regularize the NIR. Denoting the derivatives of the video signal as:

$$D(f, \theta, x, y, t) = \left( \frac{\delta f_\theta(x, y, t)}{\delta x}, \frac{\delta f_\theta(x, y, t)}{\delta y}, \frac{\delta f_\theta(x, y, t)}{\delta t} \right) \quad (7)$$

And the optical flow as:

we can now define the optical flow regularization loss

$$\mathcal{L}_{of} = \frac{1}{HWT} \sum_{x \in W} \sum_{y \in H} \sum_{t \in T} |D(f, \theta, x, y, t) \cdot F(x, y, t)| \quad (8)$$

This loss constrains the derivatives of the signal to be orthogonal to the optical flow and can be intuitively understood as keeping constant brightness along the optical flow trajectories. The total loss we use to optimize the NIR is a weighted sum of these two terms:

$$\mathcal{L} = \lambda \mathcal{L}_{obs} + (1 - \lambda) \mathcal{L}_{of} \quad (9)$$

where  $\lambda$  is a hyperparameter taking values between 0 and 1 whose impact we investigate in the following section. The exactitude of the optical flow constraint at the infinitesimal scale plays in our favor: As we regularize the true derivative of the signal representation, we do not assume constant derivatives of the signal on any interval. We believe this is the main factor behind our positive results. On the other hand, the optical flow we used was estimated from discrete consecutive frames, and thus does not represent the true infinitesimal motion field but an estimation of finite differences. We discuss this limitation in Section XXX.




Figure 3: NEED TO SHOW TRAINING CURVE AND THRESHOLD ATTAINED BY OF

## 174 4 Experiments

175 Following previous works, we use the A, B and C dataset as benchmarks to compare to the state-of-  
 176 the-art. We run all additional experiments on the XXX video illustarted in Figure XXX. Due to the  
 177 time-consuming operation of optimizing SIREN representations, we optimize and evaluate all models  
 178 on a XXX resolution. For the A dataset, we follow the standard 8 video split XXX.

179 Unless specified other-wise, all experiments are run with a SIREN of depth XXX and width XXX.  
 180 We use an omega of XXX and a lambda of XXX. We optimize the models using the Adam optimizer  
 181 using a cosine learning rate with maximum learning rate of XXX during XXX epochs.

182 We start by showing the impact of controlling the fit to high frequency without the optical flow loss in  
 183 section XXX. We show that while limiting the frequency fitted does improve generlization, it does  
 184 not allow to reach the same accuracy as optical flow regularization, showing that OF regularization  
 185 does more than just limiting the fitted frequencies.

186 In Section XXX, we compare our results to state of the art quantitatively on standard benchmarks.  
 187 We show that our approach achieves state-of-the-art resuklts on low-range motion datasets, but  
 188 underperforms existing methods on the high-range motion dataset. We present an ablation in Section  
 189 XXX, providing inisight and appropriate settings on the different model hyperparameters and a  
 190 qualitative analysis of our results in Section XXX.

191 Finally, we report a surprising additional result in Section XXX: We show that XXX.

### 192 4.1 Optical Flow constaint and High Frequencies

193 Figure 2 illustrates the fact that applying the optical flow constraint smoothes out the high-frequency  
 194 variations from the intermediate frames of vanilla SIREN representations. We start by questioning  
 195 wether the OF constraint does more than simply removing the high frequency variations of the  
 196 representation. To do so, we compare the results of vanilla SIREN representations geared towards  
 197 different frequency and compare the best obtained results to OF-constrained representations. We  
 198 constrain the SIREN frequency by varying their  $\omega$  parameter, and report our comparison in Figure  
 199 XXX.

200 While constraining the high frequency with low  $\omega$  does improve the ability to interpolate intermediate  
 201 frames, vanilla SIREN models remain well under the OF-constrained representations, confirming  
 202 than the OF constraint provides more simply restricting the high temporal frequencies.

### 203 4.2 State of the art models

204 Table XXX quantitatively compare the results of our model to state-of the art VFI models on different  
 205 datasets. We show that

Table 1: Quantitative comparison to state-of-the-art VFI on limited motion range benchmarks. Results are formatted as PSNR / SSIM.

	Adobe-240FPS [XXX]	X4K [XXX]
Super-SloMo [XXX]	27.77/0.8866	27.38/0.8527
RRIN [XXX]	32.37 / 0.9624	30.70 / 0.9270
BMBC [XXX]	27.83 / 0.9172	27.42 / 0.8585
AdaCof [XXX]	35.50 / 0.9684	34.61 / 0.9218
ABME [XXX]	35.28 / 0.9669	34.30 / 0.9195
FILM [XXX]	35.97 / 0.9710	<b>35.14</b> / 0.9397
Ours	<b>36.52 / 0.9770</b>	35.06/ <b>0.9441</b>

Table 2: Quantitative comparison to state-of-the-art VFI on large motion range benchmarks. Results are formatted as PSNR / SSIM.

	ND Scene [XXX]
V-NF	23.30 / 0.7260
NSFF [10]	28.03 / 0.9250
CURE [11]	<b>36.91 / 0.9843</b>
Ours	29.22 / 0.9215

### 206 4.3 Ablation study

207 Next, we highlight the

### 208 4.4 Qualitative Analysis

209 Ask Sho-kun.

### 210 4.5 Video fitting

## 211 5 Limitations

212 While we believe our results to be very encouraging, the proposed approach is not yet practical. Here,  
213 we discuss what we believe to be the three main limitations of, and possible solutions to, our approach

214 **Slow optimization process.** Fitting XXX frames of a video at XXX resolution currently takes XXX  
215 hours on a XXX GPU using Pytorch. This computation time is a huge draw back as it limits our  
216 ability to process full resolution video as well as to explore different hyper parameters and variants of  
217 the methods. We expect new methods speeding up the convergence of video NIR to be very benefic  
218 to this line of research. Given recent successes of NIR approaches to high impact applications (i.e.,  
219 video compression [CITE]), We hopefully expect to see advances in NIR optimisation research.

220 **Reliance on trained optical flow model.** SIREN models allow us to apply the optical flow on the  
221 exact derivaties of the signal, thus bypassing the heuristics of classical approach without relying on  
222 machine learning. The optical flow we use, however, is given by a trained ML model, which raises  
223 two problems: it is subject to generalization error, and the flow is computed on discrete samples and  
224 then subject to undesirable changes in illumination and occlusion. Future work will aim to bypass  
225 our reliance on ML-based OF using proxy constraints on the exact derivatives.

226 **Inability to interpolate high motion range videos.** In its current form, our approach only regu-  
227 larizes observed frames of the video. This has proven sufficient to reach state-of-the art on low  
228 motion ranges but is not sufficient for large motions. For larger motions several improvements can be  
229 considered, most notably by regularizing intermediate frames. Texture conservation in intermediate  
230 frames, interpolated optical flows.

## 231 6 Conclusion

232 In this paper, we have shown that regularizing NIR using the optical flow constraint equation enabled  
233 VFI without relying on ML to perform the interpolation step. We show that this approach is sufficient  
234 to reach state-of-the-art interpolation on low motion ranges

## 235 References

- 236 [1] Alexander, J.A. & Mozer, M.C. (1995) Template-based algorithms for connectionist rule extraction. In  
237 G. Tesauro, D.S. Touretzky and T.K. Leen (eds.), *Advances in Neural Information Processing Systems 7*, pp.  
238 609–616. Cambridge, MA: MIT Press.
- 239 [2] Bower, J.M. & Beeman, D. (1995) *The Book of GENESIS: Exploring Realistic Neural Models with the*  
240 *GENeral NEural Simulation System*. New York: TELOS/Springer-Verlag.
- 241 [3] Hasselmo, M.E., Schnell, E. & Barkai, E. (1995) Dynamics of learning and recall at excitatory recurrent  
242 synapses and cholinergic modulation in rat hippocampal region CA3. *Journal of Neuroscience* **15**(7):5249-5262.